

Lecture 6: September 13

Lecturer: Vijay Garg

Scribe: Sarang Bhadssavle

Agenda and Announcements

Today's lecture covered the following topics:

- The lower bound on the number of shared locations required for mutual exclusion
- Fischer's Algorithm (A timing-based algorithm)
- Lamport's Fast Mutex Algorithm
 - Splitter Construct

Additionally, remember to watch the video seminar *Toward Extreme-Scale Manycore Architectures* and the video lectures on Semaphores from last week, as they will be covered on the first exam.

6.1 Lower Bound on Number of Shared Locations

Question: Is there any algorithm which, by using just one shared variable, we can achieve mutual exclusion?

Definition 6.1 *Covering State*

The notion of a state in which all shared variables are about to be overwritten by processes and the shared state is consistent with no process in the critical section.

Theorem 6.2 (Barnes and Lynch) *Any mutex algorithm that only uses read-write variables on n processes requires at least n shared locations.*

Proof: ($n = 2$)

Let P and Q be processes and let A be a shared memory location whose initial value is \perp . Let Q run until it is about to write to A, i.e. it is in the Covering State. Now let P run and enter the critical section. Let Q run again. Q writes to A and enters the critical section.

\therefore Two processes are simultaneously in the critical section

\therefore Violation of mutual exclusion ■

This proof can be extended to $n > 2$ by following a similar sequence of events to achieve a Covering State, and then letting processes run and enter the critical section simultaneously, causing a violation of mutual exclusion.

6.2 Fischer's Algorithm

It is possible to achieve mutual exclusion of n processes with a single shared variable, provided that a key assumption is made about the timing of the algorithm. The algorithm along with this assumption is shown below.

Algorithm 1 Fischer's Algorithm (for a process P_i)

```

1: shared var turn = -1; // door is open

2: procedure REQUESTCS
3:   while(true):
4:     while(turn != -1):
5:       noOp();
6:       turn = i;
7:       wait for  $\Delta t$  time units;
8:       if(turn == i):
9:         return;

10: procedure RELEASECS
11:   turn = -1;
```

The key assumption with Fischer's Algorithm is that $\Delta t \geq c$, where c is the maximum time required to close the door i.e. the time required to set *turn* to i .

Without this assumption, the following sequence of events could occur and cause a violation of mutual exclusion for two processes, P_i and P_j :

1. P_i reads *turn* = -1
2. P_j reads *turn* = -1
3. P_j sets *turn* = j
4. P_j reads *turn* = j , and enters critical section
5. P_i sets *turn* = i
6. P_i reads *turn* = i and enters critical section

Fischer's algorithm satisfies mutex by ensuring that the time period that a process P_j waits after setting *turn* = j and before checking if *turn* == j is greater than the time period that another process P_i takes after reading *turn* == -1 and before setting *turn* = i . This ensures that only one process (in this case, P_i) can enter the critical section.

6.3 Lamport's Fast Mutex Algorithm

Lamport's Fast Mutex Algorithm is used to provide mutual exclusion in a very fast way when there is no contention for the critical section. When there is contention, the regular method of mutual exclusion is followed.

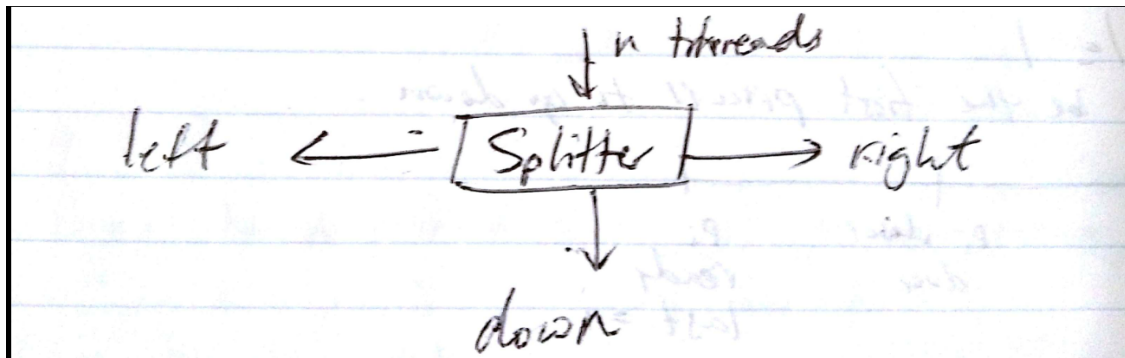


Figure 6.1: A Splitter Construct

Splitter

The construct of a "splitter" is used in Lamport's Fast Mutex algorithm, as seen in Figure 6.1.

The Splitter guarantees that:

- The number of processes sent right $\leq n - 1$
- The number of processes sent left $\leq n - 1$
- The number of processes sent down ≤ 1

Algorithm 2 Splitter Construct (for a process P_i)

```

1: shared var door : {open, closed} init open;
2: shared var last : pid init -1;

3: last = i;
4: if(door == closed):
5:   return left;
6: else:
7:   door = closed;
8:   if(last == i): return down;
9:   else: return right;
```

Claim 6.3 $\|left\| \leq n - 1$

Proof: Some process *must* have closed the door. ■

Claim 6.4 $\|right\| \leq n - 1$

Proof: Consider the process P_i such that $last = i$. Then P_i must be part of left or down. ■

Claim 6.5 $\|down\| \leq 1$

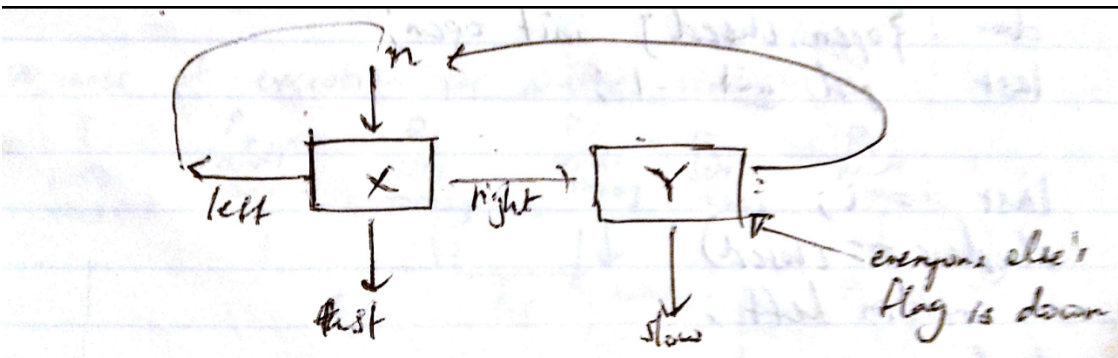


Figure 6.2: Design of Lamport's Fast Mutex Algorithm using Splitter

Proof: Let a process P_i be the first process to go down. In order to go down, P_i must execute 3 steps:

1. P_i writes *last*
2. P_i closes *door*
3. P_i reads $last == i$

Suppose there is some other process P_j also in the splitter. P_i must write *last* at some point. There are 3 cases:

- Suppose P_j writes *last* before 1. Then P_j would be the first process to go down, which violates the initial supposition that P_i went down first
- Suppose P_j writes *last* after 1 but before 3. However, this is not possible because P_i read $last == i$.
- Suppose P_j writes *last* after 3. However, this is not possible because P_j could not have found *door* to be open in the first place.

∴ Since all 3 cases are shown to be impossible, there can only be at most one process that goes down in the Splitter. ■

Lamport's Fast Mutex algorithm can be constructed using Splitter, as seen in Figure 6.2.

Proposition 6.6 Note that Lamport's Fast Mutex Algorithm does not guarantee starvation-freedom

Algorithm 3 Lamport's Fast Mutex Algorithm (for a process P_i)

```
1: procedure ACQUIRE
2:   while(true):
3:     flag[i] = up;
4:     x = i;
5:     if(y != i): // splitter's left
6:       flag[i] = down;
7:       waitUntil(y == -1);
8:       continue;
9:     else:
10:      y = i;
11:      if(x == i): return; // fast path
12:      else: // splitter's right
13:        flag[i] = down;
14:        waitUntil( $\forall j : \textit{flag}[j] == \textit{down}$ );
15:        if(y == i): return; // slow path
16:        else:
17:          waitUntil(y == -1);
18:          continue;

19: procedure RELEASE
20:   y = -1;
21:   flag[i] = down;
```

References

- [1] VIJAY GARG. EE382C, Multicore Computing, The University of Texas at Austin. September 2016.